



Implementing STM

Misha Kozik
[@mishadoff](https://twitter.com/mishadoff)

What is STM?

Software Transactional Memory

just another one
concurrency control mechanism
like locks or actors

What is STM?

Software Transactional Memory

Algorithm for Database-like concurrency

just another one
concurrency control mechanism
like locks or actors

Database Transactions

Atomicity

“all or nothing”

Consistency

data is valid at the end

Isolation

not affected by other txs

Durability

~persistence

~~Memory~~ Database Transactions

Atomicity

“all or nothing”

Consistency

data is valid at the end

Isolation

not affected by other txs

Durability

~~- persistence~~

Bank Accounts Transfer

```
public class Account {  
    private long money;  
  
    void add(long amount) {  
        money += amount;  
    }  
}  
  
public class Bank {  
    private Account[] accounts;  
  
    void transfer(Account a1, Account a2, int amount) {  
        a1.add(-amount);  
        a2.add(amount);  
    }  
  
    public long sum() {  
        long sum = 0;  
        for (Account a : accounts) sum += a.getMoney();  
    }  
}
```

Simulate Process

```
public void simulate(int threads, int num) {  
    ExecutorService service =  
        Executors.newFixedThreadPool(threads);  
    for (int i = 0; i < threads; i++) {  
        service.submit(new BankThread(this, num));  
    }  
    service.shutdown();  
    service.awaitTermination(1, TimeUnit.MINUTES);  
}  
  
public class BankThread implements Runnable {  
    private Bank bank;  
    private int num;  
  
    public void run() {  
        for (int i = 0; i < num; i++) {  
            bank.transfer(bank.getRandomAccount(),  
                          bank.getRandomAccount(),  
                          bank.getRandomValue());  
        }  
    }  
}
```

Where is my money?!

```
Bank bank = new Bank();
System.out.println("Bank sum before: " + bank.sum());
bank.simulate(100, 100000);
System.out.println("Bank sum after: " + bank.sum());
```

Bank sum before:	10000000
Bank sum after:	9970464

- 20000\$ = 
49634

How to solve it?



synchronized



What the difference?

```
public class Bank {  
    synchronized void transfer(Account a1, Account a2, int amt) {  
        a1.add(-amt);  
        a2.add(amt);  
    }  
}
```

VS

```
public class Account {  
    synchronized void transfer(Account a, int amt) {  
        this.add(-amt);  
        a.add(amt);  
    }  
  
    synchronized void add(int amt) {  
        value += amt;  
    }  
}
```

synchronized: cons

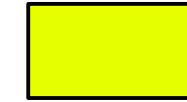
- no read/write distinction
- slow/error-prone
- low-level
- protect code rather than data

Locks



still use locks in 2013

```
RingBuffer buffer = null;
readCpsLock.lock();
try {
    if (ticks.containsKey(accountId)) {
        buffer = ticks.get(accountId);
        // check if this value changed
        if (buffer.getMaxCallCount() != cpsValue) {
            readCpsLock.unlock();
            try {
                writeCpsLock.lock();
                try {
                    ticks.remove(accountId);
                    buffer = new RingBuffer(cpsValue, DEFAULT_TIME_UNIT);
                    ticks.put(accountId, buffer);
                }
                finally {
                    writeCpsLock.unlock();
                }
            }
            finally {
                readCpsLock.lock();
            }
        }
    }
    else {
        buffer = new RingBuffer(cpsValue, DEFAULT_TIME_UNIT);
        readCpsLock.unlock();
        try {
            writeCpsLock.lock();
            try {
                ticks.put(accountId, buffer);
            }
            finally {
                writeCpsLock.unlock();
            }
        }
        finally {
            readCpsLock.lock();
        }
    }
}
finally {
    readCpsLock.unlock();
}
```



useful code



Locks: cons

- tricky
- error-prone
- not composable
- protect code rather than data

Actors



Actors in Erlang



```
ping(0, PID) ->  
    PID ! finished,  
    io:format("Ping stop~n", []).
```

```
ping(N, PID) ->  
    PID ! {ping, self()},  
    receive  
        pong ->  
            io:format("Ping~n", [])  
    end,  
    ping(N - 1, PID).
```

```
pong() ->  
    receive  
        finished ->  
            io:format("Pong stop~n", []);  
        {ping, PID} ->  
            io:format("Pong~n", []),  
            PID ! pong,  
            pong()  
    end.
```

```
start() ->  
    PID = spawn(pingpong, pong, []),  
    spawn(pingpong, ping, [3, PID]).
```

Actors: cons

- code restructuring
- deadlocks still possible

STM



ForGIFS.com

STM in Clojure



```
(def account (ref 100))
```

```
(defn transfer [acc1 acc2 amount]
  (dosync
    (alter acc1 - amount)
    (alter acc2 + amount)))
```

STM in Java

```
transaction {  
    acc1.add(-value);  
    acc2.add(value);  
}
```



STM in Java

```
STM.transaction(  
    new Runnable() {  
        public void run() {  
            acc1.add(-value);  
            acc2.add(value);  
        }  
    } );
```

STM in Java 8

```
STM.transaction(() -> {  
    acc1.add(-value);  
    acc2.add(value);  
});
```

STM: pros

- natural approach
- coordinated state
- no deadlocks, no livelocks



Implementing STM

Intuition

- start transaction
- get ref snapshot
- set of reads/writes* on snapshot
- CAS initial and current snapshot
- if CAS succeeds commit
- else retry

Ref v0.1

```
public final class Ref<T> {  
    private T value;  
  
    public T getValue() {  
        return value;  
    }  
  
    public void setValue(T value) {  
        this.value = value;  
    }  
}
```

Ref v0.1

```
public final class Ref<T> {  
    private T value;  
  
    public T getValue() {  
        return value;  
    } // from anywhere  
  
    public void setValue(T value) {  
        this.value = value;  
    } // from transaction  
}
```

Context

`<T> T get(Ref<T> ref);`

Transaction

GlobalContext

`<T> void set(Ref<T> ref, T value);`

singleton

Ref v0.2

```
public final class Ref<T> {
    private T value;

    public T getValue(Context ctx) {
        return ctx.get(this);
    }

    public void setValue(T value, Transaction tx) {
        tx.set(this, value);
    }
}
```

STM API

```
public final class STM {  
    private STM() {}  
  
    public static void transaction(TransactionBlock block) {  
        Transaction tx = new Transaction();  
        block.setTx(tx);  
        block.run();  
    }  
  
}  
  
public abstract class TransactionBlock implements Runnable {  
    private Transaction tx;  
  
    void setTx(Transaction tx) {  
        this.tx = tx;  
    }  
  
    public Transaction getTx() {  
        return tx;  
    }  
}
```

STM API

```
public final class STM {  
    private STM() {}  
  
    public static void transaction(TransactionBlock block) {  
        Transaction tx = new Transaction();  
        block.setTx(tx);  
        block.run();  
    }  
}  
  
public abstract class TransactionBlock implements Runnable {  
    private Transaction tx;  
  
    void setTx(Transaction tx) {  
        this.tx = tx;  
    }  
  
    public Transaction getTx() {  
        return tx;  
    }  
}
```



NO return value (1)

STM Usage

```
STM.transaction(new TransactionBlock() {

    @Override
    public void run() {
        int old1 = a1.getValue(this.getTx());
        a1.setValue(old1 - value, this.getTx());
        int old2 = a2.getValue(this.getTx());
        a2.setValue(old2 + value, this.getTx());
    }

});
```

STM Usage

```
STM.transaction(new TransactionBlock() {  
  
    @Override  
    public void run() {  
        int old1 = a1.getValue(this.getTx());  
        a1.setValue(old1 - value, this.getTx());  
        int old2 = a2.getValue(this.getTx());  
        a2.setValue(old2 + value, this.getTx());  
    }  
});
```

What the mess? (2)

GlobalContext v0.1

```
public class GlobalContext extends Context {  
    private HashMap<Ref, Object> refs;  
  
    @Override  
    <T> T get(Ref<T> ref) {  
        return (T)refs.get(ref);  
    }  
}
```

GlobalContext v0.1

```
public class GlobalContext extends Context {  
    private HashMap<Ref, Object> refs;  
  
    @Override  
    <T> T get(Ref<T> ref) {  
        return (T)refs.get(ref);  
    }  
}
```

need control over ref map
cheat (3)

Ref v0.3

```
public final class Ref<T> {  
    private T value;  
  
    public T getValue(Context ctx) {  
        return ctx.get(this);  
    }  
  
    public void setValue(T value, Transaction tx) {  
        tx.set(this, value);  
    }  
}
```

GlobalContext v0.2

```
public class GlobalContext extends Context {  
    private HashMap<Ref, Object> refs;  
  
    @Override  
    <T> T get(Ref<T> ref) {  
        return ref.value;  
    }  
}
```

Transaction v0.1

```
public final class Transaction extends Context {  
    private HashMap<Ref, Object> inTxMap;  
  
    @Override  
    <T> T get(Ref<T> ref) {  
        if (!inTxMap.containsKey(ref)) {  
            inTxMap.put(ref, ref.value);  
        }  
        return (T)inTxMap.get(ref);  
    }  
  
    <T> void set(Ref<T> ref, T value) {  
        inTxMap.put(ref, value);  
    }  
}
```

Transaction v0.1

```
public final class Transaction extends Context {  
    private HashMap<Ref, Object> inTxMap;
```

```
@Override
```

NO Snapshot Isolation (4)

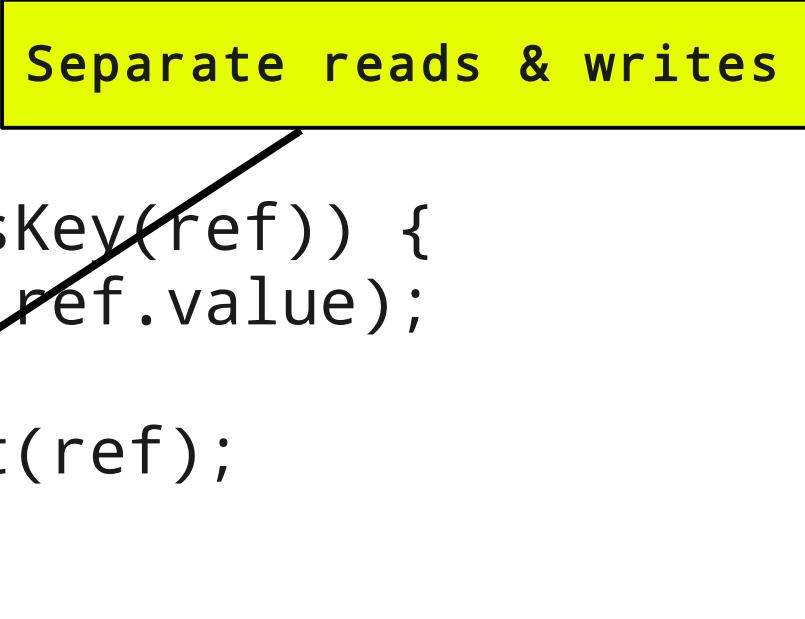
```
<T> T get(Ref<T> ref) {  
    if (!inTxMap.containsKey(ref)) {  
        inTxMap.put(ref, ref.value);  
    }  
    return (T)inTxMap.get(ref);  
}
```

```
<T> void set(Ref<T> ref, T value) {  
    inTxMap.put(ref, value);  
}
```

```
}
```

Transaction v0.1

```
public final class Transaction extends Context {  
    private HashMap<Ref, Object> inTxMap;  
  
    @Override  
    <T> T get(Ref<T> ref) {  
        if (!inTxMap.containsKey(ref)) {  
            inTxMap.put(ref, ref.value);  
        }  
        return (T)inTxMap.get(ref);  
    }  
  
    <T> void set(Ref<T> ref, T value) {  
        inTxMap.put(ref, value);  
    }  
}
```



Separate reads & writes

Transaction v0.2

```
public final class Transaction extends Context {  
    private HashMap<Ref, Object> inTxMap;  
    private HashSet<Ref> toUpdate;  
  
    @Override  
    <T> T get(Ref<T> ref) {  
        if (!inTxMap.containsKey(ref)) {  
            inTxMap.put(ref, ref.value);  
        }  
        return (T)inTxMap.get(ref);  
    }  
  
    <T> void set(Ref<T> ref, T value) {  
        inTxMap.put(ref, value);  
        toUpdate.add(ref);  
    }  
}
```

Commit

```
public final class Transaction extends Context {  
    // ...  
  
    void commit() {  
        synchronized (STM.commitLock) {  
            // TODO validate  
            // TODO write values  
        }  
    }  
  
}  
    public final class STM {  
        private STM() {}  
  
        public static Object commitLock = new Object();  
  
        // ...  
    }
```

Commit: Write Values

```
void commit() {  
    synchronized (STM.commitLock) {  
        // TODO validate  
        for (Ref ref : toUpdate) {  
            ref.value = inTxMap.get(ref);  
        }  
    }  
}
```

Validation Idea

- get all transaction-local refs
- compare version with transaction revision
- if all versions the same
allow commit
- else retry

Ref v0.4

```
public final class Ref<T> {  
    T value;  
    long revision = 0;  
  
    // ...  
}
```

Ref v0.4

```
public final class Ref<T> {  
    T value;  
    long revision = 0;  
    // ...  
}
```

No history (5)

Transaction v0.3

```
public final class Transaction extends Context {  
    // ...  
  
    private long revision;  
    private static AtomicLong transactionNum =  
        new AtomicLong(0);  
  
    Transaction() {  
        revision = transactionNum.incrementAndGet();  
    }  
  
    // ...  
}
```

Transaction v0.3

```
public final class Transaction extends Context {  
    private HashMap<Ref, Long> version;  
  
    <T> T get(Ref<T> ref) {  
        if (!inTxMap.containsKey(ref)) {  
            inTxMap.put(ref, ref.value);  
            if (!version.containsKey(ref)) {  
                version.put(ref, ref.revision);  
            }  
        }  
        return (T)inTxMap.get(ref);  
    }  
  
    <T> void set(Ref<T> ref, T value) {  
        inTxMap.put(ref, value);  
        toUpdate.add(ref);  
        if (!version.containsKey(ref)) {  
            version.put(ref, ref.revision);  
        }  
    }  
}
```

Commit: Validation

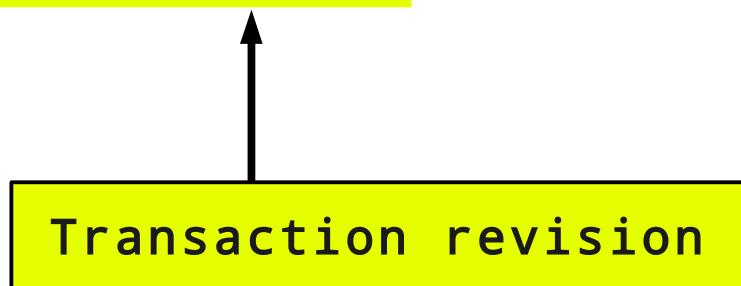
```
boolean commit() {  
    synchronized (STM.commitLock) {  
        boolean isValid = true;  
        for (Ref ref : inTxMap.keySet()) {  
            if (ref.revision != version.get(ref)) {  
                isValid = false;  
                break;  
            }  
        }  
  
        if (isValid) {  
            for (Ref ref : toUpdate) {  
                ref.value = inTxMap.get(ref);  
                ref.revision = revision;  
            }  
        }  
        return isValid;  
    }  
}
```

Commit: Validation

```
boolean commit() {  
    synchronized (STM.commitLock) {  
        boolean isValid = true;  
        for (Ref ref : inTxMap.keySet()) {  
            if (ref.revision != version.get(ref)) {  
                isValid = false;  
                break;  
            }  
        }  
  
        if (isValid) {  
            for (Ref ref : toUpdate) {  
                ref.value = inTxMap.get(ref);  
                ref.revision = revision;  
            }  
        }  
        return isValid;  
    }  
}
```

Commit: Validation

```
boolean commit() {  
    synchronized (STM.commitLock) {  
        boolean isValid = true;  
        for (Ref ref : inTxMap.keySet()) {  
            if (ref.revision != version.get(ref)) {  
                isValid = false;  
                break;  
            }  
        }  
  
        if (isValid) {  
            for (Ref ref : toUpdate) {  
                ref.value = inTxMap.get(ref);  
                ref.revision = revision;  
            }  
        }  
        return isValid;  
    }  
}
```



Commit: Validation

```
boolean commit() {  
    synchronized (STM.commitLock) {  
        boolean isValid = true;  
        for (Ref ref : inTxMap.keySet()) {  
            if (ref.revision != version.get(ref)) {  
                isValid = false;  
                break;  
            }  
        }  
  
        if (isValid) {  
            for (Ref ref : toUpdate) {  
                ref.value = inTxMap.get(ref);  
                ref.revision = revision;  
            }  
        }  
        return isValid;  
    }  
}
```

STM Transaction

```
public final class STM {  
    private STM() {}  
  
    public static Object commitLock = new Object();  
  
    public static void transaction(TransactionBlock block) {  
        boolean committed = false;  
        while (!committed) {  
            Transaction tx = new Transaction();  
            block.setTx(tx);  
            block.run();  
            committed = tx.commit();  
        }  
    }  
}
```

STM Transaction

```
public final class STM {  
    private STM() {}  
  
    public static Object commitLock = new Object();  
  
    public static void transaction(TransactionBlock block) {  
        boolean committed = false;  
        while (!committed) {  
            Transaction tx = new Transaction();  
            block.setTx(tx);  
            block.run();  
            committed = tx.commit();  
        }  
    }  
}
```

No exceptions
handling (6)

No nested
transactions (7)

BRO...

IS IT WORKING?

What the problem?

```
boolean commit() {  
    synchronized (STM.commitLock) {  
        // ...  
        if (isValid) {  
            for (Ref ref : toUpdate) {  
                ref.value = inTxMap.get(ref);  
                ref.revision = revision;  
            }  
        }  
        // ...  
    }  
}
```

What the problem?

```
boolean commit() {  
    synchronized (STM.commitLock) {  
        // ...  
        if (isValid) {  
            for (Ref ref : toUpdate) {  
                ref.value = inTxMap.get(ref);  
                ref.revision = revision;  
            }  
        }  
        // ...  
    }  
}
```

Memory switch is atomic

Two Memory switches
are not atomic

RefTuple

```
public class RefTuple<V, R> {
    V value;
    R revision;

    public RefTuple(V v, R r) {
        this.value = v;
        this.revision = r;
    }

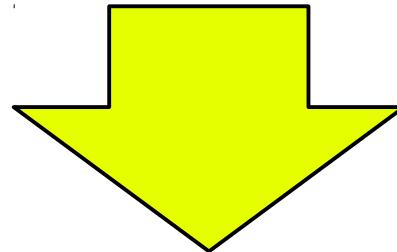
    static <V, R> RefTuple get(V v, R r) {
        return new RefTuple<V, R>(v, r);
    }
}
```

Ref v0.5

```
public final class Ref<T> {  
    RefTuple<T, Long> content;  
  
    public Ref(T value) {  
        content = RefTuple.get(value, 0);  
    }  
  
    // ...  
}
```

Transaction.commit()

```
for (Ref ref : toUpdate) {  
    ref.value = inTxMap.get(ref);  
    ref.revision = revision;  
}
```

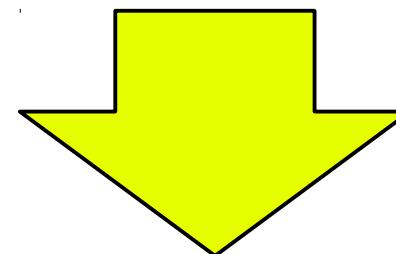


```
for (Ref ref : toUpdate) {  
    ref.content =  
        RefTuple.get(inTxMap.get(ref), revision);  
}
```

Transaction.get()

```
<T> T get(Ref<T> ref) {  
    if (!inTxMap.containsKey(ref)) {  
        inTxMap.put(ref, ref.value);  
        if (!version.containsKey(ref)) {  
            version.put(ref, ref.revision);  
        }  
    }  
    return (T)inTxMap.get(ref);  
}
```

2 non-sync reads



```
<T> T get(Ref<T> ref) {  
    if (!inTxMap.containsKey(ref)) {  
        RefTuple<T, Long> tuple = ref.content;  
        inTxMap.put(ref, tuple.value);  
        if (!version.containsKey(ref)) {  
            version.put(ref, tuple.revision);  
        }  
    }  
    return (T)inTxMap.get(ref);  
}
```

// TODO

1. Support return value in TX
2. Avoid `this.getTx()` repetition
3. Fair **GlobalContext**
4. Snapshot Isolation
5. Support ref history
6. Support exceptions
7. Support nested transactions

STM Algorithms

Multi-Version Concurrency Control

Lazy Snapshot Algorithm

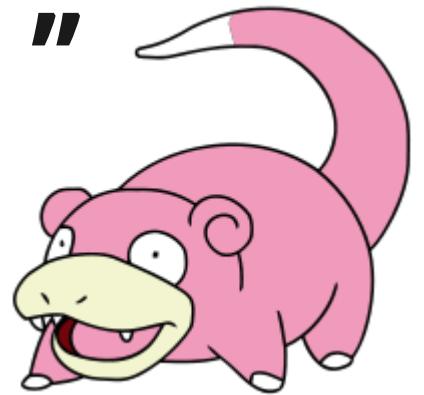
Transactional Locking 2

Practical Advices

- avoid side-effects
- avoid short/long transactions
- separate as much as possible outside of transaction
- ref value must be immutable

OR STM WON'T HELP YOU

"But STM is slow..."



No!

TM/GC Analogy

Transactional Memory is to
shared-memory concurrency

as

Garbage Collection is to
memory management.

Non-STM

- atomics
- immutability
- avoid shared state
- lock-free data structures



Links

- STM In Clojure
- Locks, Actors and STM in pictures
- The Dining Philosophers Problem
- Clojure refs
- TM/GC Analogy [PDF]
- STM in Scala
- 7 Concurrency Models in 7 weeks